


MyIntrinsics++ (MIPP)

pipeline passed coverage 99.20%



Purpose

MIPP is a portable and Open-source wrapper (MIT license) for vector intrinsic functions (SIMD) written in C++11. It works for SSE, AVX, AVX-512, ARM NEON and SVE (work in progress) instructions. MIPP wrapper supports simple/double precision floating-point numbers and also signed/unsigned integer arithmetic (64-bit, 32-bit, 16-bit and 8-bit).

With the MIPP wrapper you do not need to write a specific intrinsic code anymore. Just use provided functions and the wrapper will automatically generates the right intrinsic calls for your specific architecture.

If you are interested by ARM SVE development status, [please follow this link](#).

Short Documentation

Supported Compilers

At this time, MIPP has been tested on the following compilers:

- Intel: `icpc` ≥ 16 ,
- GNU: `g++` ≥ 4.8 ,
- Clang: `clang++` ≥ 3.6 ,
- Microsoft: `msvc` ≥ 14 .

On `msvc 14.10` (Microsoft Visual Studio 2017), the performances are reduced compared to the other compilers, the compiler is not able to fully inline all the MIPP methods. This has been fixed on `msvc 14.21` (Microsoft Visual Studio 2019) and now you can expect high performances.

Install and Configure your Code

You don't have to install MIPP because it is a simple C++ header file. The headers are located in the `include` folder (note that this location has changed since commit `6795891`, before they were located in the `src` folder).

Just include the header into your source files when the wrapper is needed.

```
#include "mipp.h"
```

mipp.h use a C++ `namespace`: `mipp`, if you do not want to prefix all the MIPP calls by `mipp::` you can do that:

```
#include "mipp.h"
using namespace mipp;
```

Before trying to compile, think to tell the compiler what kind of vector instructions you want to use. For instance, if you are using GNU compiler (`g++`) you simply have to add the `-march=native` option for SSE and AVX CPUs compatible. For ARMv7 CPUs with NEON instructions you have to add the `-mfpu=neon` option (since most of current NEONv1 instructions are not IEEE-754 compliant). However, this is no more the case on ARMv8 processors, so the `-march=native` option will work too. MIPP also uses some nice features provided by the C++11 and so we have to add the `-std=c++11` flag to compile the code. You are now ready to run your code with the MIPP wrapper.

In the case where MIPP is installed on the system it can be integrated into a cmake projet in a standard way. Example

```
# install MIPP
cd MIPP/
export MIPP_ROOT=$PWD/build/install
cmake -B build -DCMAKE_INSTALL_PREFIX=$MIPP_ROOT
cmake --build build -j5
cmake --install build
```

In your `CMakeLists.txt`:

```
# find the installation of MIPP on the system
find_package(MIPP REQUIRED)

# define your executable
add_executable(gemm gemm.cpp)

# link your executable to MIPP
target_link_libraries(gemm PRIVATE MIPP::mipp)
```

```
cd your_project/
# if MIPP is installed in a system standard path: MIPP will be found automatically
# with cmake
cmake -B build
# if MIPP is installed in a non-standard path: use CMAKE_PREFIX_PATH
cmake -B build -DCMAKE_PREFIX_PATH=$MIPP_ROOT
```

Generate Sources & Compile the Static Library

MIPP is mainly a header only library. However, some macro operations require to compile a small library. This is particularly true for the `compress` operation that relies on generated LUTs stored in the static library.

To generate the source files containing these LUTs you need to install Python3 with the Jinja2 package:

```
sudo apt install python3 python3-pip
pip3 install --user -r codegen/requirements.txt
```

Then you can call the generator as follow:

```
python3 codegen/gen_compress.py
```

And, finally you can compile the MIPP static library:

```
cmake -B build -DMIPP_STATIC_LIB=ON
cmake --build build -j4
```

Note that **the compilation of the static library is optional**. You can choose to do not compile the static library then only some macro operations will be missing.

Sequential Mode

By default, MIPP tries to recognize the instruction set from the preprocessor definitions. If MIPP can't match the instruction set (for instance when MIPP does not support the targeted instruction set), MIPP falls back on standard sequential instructions. In this mode, the vectorization is not guarantee anymore but the compiler can still perform auto-vectorization.

It is possible to force MIPP to use the sequential mode with the following compiler definition: - `DMIPP_NO_INTRINSICS`. Sometime it can be useful for debugging or to bench a code.

If you want to check the MIPP mode configuration, you can print the following global variable: `mipp::InstructionFullType` (`std::string`).

Vector Register Declaration

Just use the `mipp::Reg<T>` type.

```
mipp::Reg<T> r1, r2, r3; // we have declared 3 vector registers
```

But we do not know the number of elements per register here. This number of elements can be obtained by calling the `mipp::N<T>()` function (`T` is a template parameter, it can be `double`, `float`, `int64_t`, `uint64_t`, `int32_t`, `uint32_t`, `int16_t`, `uint16_t`, `int8_t` or `uint8_t` type).

```
for (int i = 0; i < n; i += mipp::N<float>()) {
    // ...
}
```

The register size directly depends on the precision of the data we are working on.

Register **load** and **store** Instructions

Loading memory from a vector into a register:

```
int n = mipp::N<float>() * 10;
std::vector<float> myVector(n);
int i = 0;
mipp::Reg<float> r1;
r1.load(&myVector[i*mipp::N<float>()]);
```

The last two lines can be shorten as follow where the **load** call becomes implicit:

```
mipp::Reg<float> r1 = &myVector[i*mipp::N<float>()];
```

Store can be done with the **store(...)** method:

```
int n = mipp::N<float>() * 10;
std::vector<float> myVector(n);
int i = 0;
mipp::Reg<float> r1 = &myVector[i*mipp::N<float>()];

// do something with r1

r1.store(&myVector[(i+1)*mipp::N<float>()]);
```

By default the loads and stores work on **unaligned memory**. It is possible to control this behavior with the `-DMIPP_ALIGNED_LOADS` definition: when specified, the loads and stores work on **aligned memory** by default. In the **aligned memory** mode, it is still possible to perform unaligned memory operations with the `mipp::loadu` and `mipp::storeu` functions. However, it is not possible to perform aligned loads and stores in the **unaligned memory** mode.

To allocate aligned data you can use the MIPP aligned memory allocator wrapped into the `mipp::vector` class. `mipp::vector` is fully retro-compatible with the standard `std::vector` class and it can be use everywhere you can use `std::vector`.

```
mipp::vector<float> myVector(n);
```

Register Initialization

You can initialize a vector register from a scalar value:

```
mipp::Reg<float> r1; // r1 = | unknown | unknown | unknown | unknown |
r1 = 1.0;           // r1 = |   +1.0 |   +1.0 |   +1.0 |   +1.0 |
```

Or from an initializer list (`std::initializer_list`):

```
mipp::Reg<float> r1; // r1 = | unknown | unknown | unknown | unknown |
r1 = {1.0, 2.0, 3.0, 4.0}; // r1 = |   +1.0 |   +2.0 |   +3.0 |   +4.0 |
```

Computational Instructions

Add two vector registers:

```
mipp::Reg<float> r1, r2, r3;

r1 = 1.0; // r1 = | +1.0 | +1.0 | +1.0 | +1.0 |
r2 = 2.0; // r2 = | +2.0 | +2.0 | +2.0 | +2.0 |

r3 = r1 + r2; // r3 = | +3.0 | +3.0 | +3.0 | +3.0 |
```

Subtract two vector registers:

```
mipp::Reg<float> r1, r2, r3;

r1 = 1.0; // r1 = | +1.0 | +1.0 | +1.0 | +1.0 |
r2 = 2.0; // r2 = | +2.0 | +2.0 | +2.0 | +2.0 |

r3 = r1 - r2; // r3 = | -1.0 | -1.0 | -1.0 | -1.0 |
```

Multiply two vector registers:

```
mipp::Reg<float> r1, r2, r3;

r1 = 1.0; // r1 = | +1.0 | +1.0 | +1.0 | +1.0 |
r2 = 2.0; // r2 = | +2.0 | +2.0 | +2.0 | +2.0 |

r3 = r1 * r2; // r3 = | +2.0 | +2.0 | +2.0 | +2.0 |
```

Divide two vector registers:

```
mipp::Reg<float> r1, r2, r3;

r1 = 1.0;      // r1 = | +1.0 | +1.0 | +1.0 | +1.0 |
r2 = 2.0;      // r2 = | +2.0 | +2.0 | +2.0 | +2.0 |

r3 = r1 / r2;  // r3 = | +0.5 | +0.5 | +0.5 | +0.5 |
```

Fused multiply and add of three vector registers:

```
mipp::Reg<float> r1, r2, r3, r4;

r1 = 2.0;      // r1 = | +2.0 | +2.0 | +2.0 | +2.0 |
r2 = 3.0;      // r2 = | +3.0 | +3.0 | +3.0 | +3.0 |
r3 = 1.0;      // r3 = | +1.0 | +1.0 | +1.0 | +1.0 |

// r4 = (r1 * r2) + r3
r4 = mipp::fmadd(r1, r2, r3); // r4 = | +7.0 | +7.0 | +7.0 | +7.0 |
```

Fused negative multiply and add of three vector registers:

```
mipp::Reg<float> r1, r2, r3, r4;

r1 = 2.0;      // r1 = | +2.0 | +2.0 | +2.0 | +2.0 |
r2 = 3.0;      // r2 = | +3.0 | +3.0 | +3.0 | +3.0 |
r3 = 1.0;      // r3 = | +1.0 | +1.0 | +1.0 | +1.0 |

// r4 = -(r1 * r2) + r3
r4 = mipp::fnmadd(r1, r2, r3); // r4 = | -5.0 | -5.0 | -5.0 | -5.0 |
```

Square root of a vector register:

```
mipp::Reg<float> r1, r2;

r1 = 9.0;      // r1 = | +9.0 | +9.0 | +9.0 | +9.0 |

r2 = mipp::sqrt(r1); // r2 = | +3.0 | +3.0 | +3.0 | +3.0 |
```

Reciprocal square root of a vector register (be careful: this intrinsic exists only for simple precision floating-point numbers):

```
mipp::Reg<float> r1, r2;

r1 = 9.0;      // r1 = | +9.0 | +9.0 | +9.0 | +9.0 |
```

```
r2 = mipp::rsqrt(r1); // r2 = | +0.3 | +0.3 | +0.3 | +0.3 |
```

Selections

Select the **minimum** between two vector registers:

```
mipp::Reg<float> r1, r2, r3;

r1 = 2.0;           // r1 = | +2.0 | +2.0 | +2.0 | +2.0 |
r2 = 3.0;           // r2 = | +3.0 | +3.0 | +3.0 | +3.0 |

r3 = mipp::min(r1, r2); // r3 = | +2.0 | +2.0 | +2.0 | +2.0 |
```

Select the **maximum** between two vector registers:

```
mipp::Reg<float> r1, r2, r3;

r1 = 2.0;           // r1 = | +2.0 | +2.0 | +2.0 | +2.0 |
r2 = 3.0;           // r2 = | +3.0 | +3.0 | +3.0 | +3.0 |

r3 = mipp::max(r1, r2); // r3 = | +3.0 | +3.0 | +3.0 | +3.0 |
```

Permutations

The `rrot(...)` method allows you to perform a **right rotation** (a cyclic permutation) of the elements inside the register:

```
mipp::Reg<float> r1, r2;
r1 = {3.0, 2.0, 1.0, 0.0} // r1 = | +3.0 | +2.0 | +1.0 | +0.0 |

r2 = mipp::rrot(r1);      // r2 = | +0.0 | +3.0 | +2.0 | +1.0 |
r1 = mipp::rrot(r2);      // r1 = | +1.0 | +0.0 | +3.0 | +2.0 |
r2 = mipp::rrot(r1);      // r2 = | +2.0 | +1.0 | +0.0 | +3.0 |
r1 = mipp::rrot(r2);      // r1 = | +3.0 | +2.0 | +1.0 | +0.0 |
```

Of course there are many more available instructions in the MIPP wrapper and you can find these instructions at the [end of this page](#).

Addition of Two Vectors

```
#include <cstdlib> // rand()
#include "mipp.h"
```

```

int main()
{
    // data allocation
    const int n = 32000; // size of the vA, vB, vC vectors
    mipp::vector<float> vA(n); // in
    mipp::vector<float> vB(n); // in
    mipp::vector<float> vC(n); // out

    // data initialization
    for (int i = 0; i < n; i++) vA[i] = rand() % 10;
    for (int i = 0; i < n; i++) vB[i] = rand() % 10;

    // declare 3 vector registers
    mipp::Reg<float> rA, rB, rC;

    // compute rC with the MIPP vectorized functions
    for (int i = 0; i < n; i += mipp::N<float>()) {
        rA.load(&vA[i]); // unaligned load by default (use the -
DMIPP_ALIGNED_LOADS
        rB.load(&vB[i]); // macro definition to force aligned loads and stores).
        rC = rA + rB;
        rC.store(&vC[i]);
    }

    return 0;
}

```

Vectorizing an Existing Code

Scalar Code

```

// ...
for (int i = 0; i < n; i++) {
    out[i] = 0.75f * in1[i] * std::exp(in2[i]);
}
// ...

```

Vectorized Code

```

// ...
// compute the vectorized loop size which is a multiple of 'mipp::N<float>()'.
auto vecLoopSize = (n / mipp::N<float>()) * mipp::N<float>();
mipp::Reg<float> rout, rin1, rin2;
for (int i = 0; i < vecLoopSize; i += mipp::N<float>()) {
    rin1.load(&in1[i]); // unaligned load by default (use the -DMIPP_ALIGNED_LOADS
    rin2.load(&in2[i]); // macro definition to force aligned loads and stores).
    // the '0.75f' constant will be broadcast in a vector but it has to be at
    // the right of a 'mipp::Reg<T>', this is why it has been moved at the right

```



```

    // of the 'rin1' register. Notice that 'std::exp' has been replaced by
    // 'mipp::exp'.
    rout = rin1 * 0.75f * mipp::exp(rin2);
    rout.store(&out[i]);
}

// scalar tail loop: compute the remaining elements that can't be vectorized.
for (int i = vecLoopSize; i < n; i++) {
    out[i] = 0.75f * in1[i] * std::exp(in2[i]);
}
// ...

```

Masked Instructions

MIPP comes with two generic and templated masked functions (`mask` and `maskz`). Those functions allow you to benefit from the AVX-512 and SVE masked instructions. `mask` and `maskz` functions are retro compatible with older instruction sets.

```

mipp::Reg<        float    > ZMM1 = {  40,  -30,   60,   80};
mipp::Reg<        float    > ZMM2 = 0.1; // broadcast
mipp::Msk<mipp::N<float>()> k1   = {false, true, false, false};

// ZMM3 = k1 ? ZMM1 * ZMM2 : ZMM1;
auto ZMM3 = mipp::mask<float, mipp::mul>(k1, ZMM1, ZMM1, ZMM2);
std::cout << ZMM3 << std::endl; // output: "[40, -3, 60, 80]"

// ZMM4 = k1 ? ZMM1 * ZMM2 : 0;
auto ZMM4 = mipp::maskz<float, mipp::mul>(k1, ZMM1, ZMM2);
std::cout << ZMM4 << std::endl; // output: "[0, -3, 0, 0]"

```

List of MIPP Functions

This section presents an exhaustive list of all the available functions in MIPP. Of course the MIPP wrapper does not cover all the possible intrinsics of each instruction set but it tries to give you the most important and useful ones.

In the following tables, `T`, `T1` and `T2` stand for data types (`double`, `float`, `int64_t`, `uint64_t`, `int32_t`, `uint32_t`, `int16_t`, `uint16_t`, `int8_t` or `uint8_t`). `N` stands for the number or elements in a mask or in a register. `N` is a strictly positive integer and can easily be deduced from the data type: `constexpr int N = mipp::N<T>()`. When `T` and `N` are mixed in a prototype, `N` has to satisfy the previous constraint (`N = mipp::N<T>()`).

In the documentation there are some terms that requires to be clarified:

- **register element:** a SIMD register is composed by multiple scalar elements, those elements are built-in data types (`double`, `float`, `int64_t`, ...),
- **register lane:** modern instruction sets can have multiple implicit sub parts in an entire SIMD register, those sub parts are called lanes (SSE has one lane of 128 bits, AVX has two lanes of 128 bits, AVX-512

has four lanes of 128 bits).

Memory Operations

Short name	Prototype	Documentation	Supported types
load	Reg <T> load (const T* mem)	Loads aligned data from mem to a register.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
loadu	Reg <T> loadu (const T* mem)	Loads unaligned data from mem to a register.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
store	void store (T* mem, const Reg<T> r)	Stores the r register in the mem aligned data.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
storeu	void storeu (T* mem, const Reg<T> r)	Stores the r register in the mem unaligned data.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
maskzld	Reg <T> maskzld (const Msk<N> m, const T* mem)	Loads elements according to the mask m (puts zero when the mask value is false).	double, float, int64_t, int32_t, int16_t, int8_t
maskzlds	Reg <T> maskzlds (const Msk<N> m, const T* mem)	Loads elements according to the mask m (puts zero when the mask value is false). Safe version, only reads masked elements in memory.	double, float, int64_t, int32_t, int16_t, int8_t
maskst	void maskst (const Msk<N> m, T* mem, const Reg<T> r)	Stores elements from the r register according to the mask m in the mem memory.	double, float, int64_t, int32_t, int16_t, int8_t
masksts	void masksts (const Msk<N> m, T* mem, const Reg<T> r)	Stores elements from the r register according to the mask m in the mem memory. Safe version, only writes masked elements in memory.	double, float, int64_t, int32_t, int16_t, int8_t
gather	Reg <TD, TI> gather (const TD* mem, const Reg<TI> idx)	Gathers elements from mem to a register. Selects elements according to the indices in idx.	double, float, int64_t, int32_t, int16_t, int8_t

Short name	Prototype	Documentation	Supported types
scatter	void <TD, TI> scatter (TD* mem, const Reg<TI> idx, const Reg<TD> r)	Scatters elements into mem from the r register. Writes elements at the idx indices in mem .	double, float, int64_t, int32_t, int16_t, int8_t
maskzgat	Reg <TD, TI> gather (const Msk<N> m, const TD* mem, const Reg<TI> idx)	Gathers elements from mem to a register (according to the mask m). Selects elements according to the indices in idx (puts zero when the mask value is false).	double, float, int64_t, int32_t, int16_t, int8_t
masksca	void <TD, TI> scatter (const Msk<N> m, TD* mem, const Reg<TI> idx, const Reg<TD> r)	Scatters elements into mem from the r register (according to the mask m). Writes elements at the idx indices in mem .	double, float, int64_t, int32_t, int16_t, int8_t
set	Reg <T> set (const T[N] vals)	Sets a register from the values in vals .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
set	Msk <N> set (const bool[N] bits)	Sets a mask from the bits in bits .	
set1	Reg <T> set1 (const T val)	Broadcasts val in a register.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
set1	Msk <N> set1 (const bool bit)	Broadcasts bit in a mask.	
set0	Reg <T> set0 ()	Initializes a register to zero.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
set0	Msk <N> set0 ()	Initializes a mask to false.	
get	T get (const Reg<T> r, const size_t index)	Gets a specific element from the register r at the index position.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t

Short name	Prototype	Documentation	Supported types
get	<code>T get (const Reg_2<T> r, const size_t index)</code>	Gets a specific element from the register <code>r</code> at the <code>index</code> position.	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
get	<code>bool get (const Msk<N> m, const size_t index)</code>	Gets a specific element from the register <code>m</code> at the <code>index</code> position.	
getfirst	<code>T getfirst (const Reg<T> r)</code>	Gets the first element from the register <code>r</code> .	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
getfirst	<code>T getfirst (const Reg_2<T> r)</code>	Gets the first element from the register <code>r</code> .	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
getfirst	<code>bool getfirst (const Msk<N> m)</code>	Gets the first element from the register <code>m</code> .	
low	<code>Reg_2<T> low (const Reg<T> r)</code>	Gets the low part of the <code>r</code> register.	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
high	<code>Reg_2<T> high (const Reg<T> r)</code>	Gets the high part of the <code>r</code> register.	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
combine	<code>Reg <T> combine (const Reg_2<T> r1, const Reg_2<T> r2)</code>	Combine two half registers in a full register, <code>r1</code> will be the low part and <code>r2</code> the high part.	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
combine	<code>Reg <S,T> combine (const Reg<T> r1, const Reg<T> r2)</code>	<code>S</code> elements of <code>r1</code> are shifted to the left, <code>(S - N) + N</code> elements of <code>r2</code> are shifted to the right. Shifted <code>r1</code> and <code>r2</code> are combined to give the result.	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>

Short name	Prototype	Documentation	Supported types
<code>compress</code>	<code>Reg <T> compress (const Reg<T> r1, const Msk<N> m)</code>	Pack the elements of <code>r1</code> at the beginning of the register according to the bitmask <code>m</code> (if the bit is 1 then element is picked, otherwise it is not).	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
<code>cmask</code>	<code>Reg <T> cmask (const uint32_t[N] ids)</code>	Creates a cmask from an indexes list (indexes have to be between 0 and N-1).	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
<code>cmask2</code>	<code>Reg <T> cmask2 (const uint32_t[N/2] ids)</code>	Creates a cmask2 from an indexes list (indexes have to be between 0 and (N/2)-1).	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
<code>cmask4</code>	<code>Reg <T> cmask4 (const uint32_t[N/4] ids)</code>	Creates a cmask4 from an indexes list (indexes have to be between 0 and (N/4)-1).	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
<code>shuff</code>	<code>Reg <T> shuff (const Reg<T> r, const Reg<T> cm)</code>	Shuffles the elements of <code>r</code> according to the cmask <code>cm</code> .	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
<code>shuff2</code>	<code>Reg <T> shuff2 (const Reg<T> r, const Reg<T> cm2)</code>	Shuffles the elements of <code>r</code> according to the cmask2 <code>cm2</code> (same shuffle is applied in both lanes).	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
<code>shuff4</code>	<code>Reg <T> shuff4 (const Reg<T> r, const Reg<T> cm4)</code>	Shuffles the elements of <code>r</code> according to the cmask4 <code>cm4</code> (same shuffle is applied in the four lanes).	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>
<code>interleave</code>	<code>Regx2<T> interleave (const Reg<T> r1, const Reg<T> r2)</code>	Interleaves <code>r1</code> and <code>r2</code> : [<code>r1_1</code> , <code>r2_1</code> , <code>r1_2</code> , <code>r2_2</code> , ..., <code>r1_n</code> , <code>r2_n</code>].	<code>double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t</code>

Short name	Prototype	Documentation	Supported types
deinterleave	<code>Regx2<T></code> <code>deinterleave</code> <code>(const Reg<T> r1,</code> <code>const Reg<T> r2)</code>	Reverts the previous defined interleave operation.	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>
interleave2	<code>Regx2<T></code> <code>interleave2 (const</code> <code>Reg<T> r1, const</code> <code>Reg<T> r2)</code>	Interleaves <code>r1</code> and <code>r2</code> considering two lanes.	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>
interleave4	<code>Regx2<T></code> <code>interleave4 (const</code> <code>Reg<T> r1, const</code> <code>Reg<T> r2)</code>	Interleaves <code>r1</code> and <code>r2</code> considering four lanes.	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>
interleavelo	<code>Reg <T></code> <code>interleavelo</code> <code>(const Reg<T> r1,</code> <code>const Reg<T> r2)</code>	Interleaves the low part of <code>r1</code> with the low part of <code>r2</code> .	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>
interleavelo2	<code>Reg <T></code> <code>interleavelo2</code> <code>(const Reg<T> r1,</code> <code>const Reg<T> r2)</code>	Interleaves the low part of <code>r1</code> with the low part of <code>r2</code> (considering two lanes).	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>
interleavelo4	<code>Reg <T></code> <code>interleavelo4</code> <code>(const Reg<T> r1,</code> <code>const Reg<T> r2)</code>	Interleaves the low part of <code>r1</code> with the low part of <code>r2</code> (considering four lanes).	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>
interleavehi	<code>Reg <T></code> <code>interleavehi</code> <code>(const Reg<T> r1,</code> <code>const Reg<T> r2)</code>	Interleaves the high part of <code>r1</code> with the high part of <code>r2</code> .	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>
interleavehi2	<code>Reg <T></code> <code>interleavehi2</code> <code>(const Reg<T> r1,</code> <code>const Reg<T> r2)</code>	Interleaves the high part of <code>r1</code> with the high part of <code>r2</code> (considering two lanes).	<code>double, float,</code> <code>int64_t, uint64_t,</code> <code>int32_t, uint32_t,</code> <code>int16_t, uint16_t,</code> <code>int8_t, uint8_t</code>

Short name	Prototype	Documentation	Supported types
interleavehi4	Reg <T> interleavehi4 (const Reg<T> r1, const Reg<T> r2)	Interleaves the high part of r1 with the high part of r2 (considering four lanes).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
lrot	Reg <T> lrot (const Reg<T> r)	Rotates the r register from the left (cyclic permutation).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
rrot	Reg <T> rrot (const Reg<T> r)	Rotates the r register from the right (cyclic permutation).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
blend	Reg <T> blend (const Reg<T> r1, const Reg<T> r2, const Msk<N> m)	Combines r1 and r2 register following the m mask values (m_i ? r1_i : r2_i).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
select	Reg <T> select (const Msk<N> m, const Reg<T> r1, const Reg<T> r2)	Alias for the previous blend function. Parameters order is a little bit different.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t

Bitwise Operations

The **pipe** keyword stands for the "|" binary operator.

Short name	Operator	Prototype	Documentation	Supported types
andb	& and &=	Reg<T> andb (const Reg<T> r1, const Reg<T> r2)	Computes the bitwise AND: r1 & r2 .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
andb	& and &=	Msk<N> andb (const Msk<N> m1, const Msk<N> m2)	Computes the bitwise AND: m1 & m2 .	

Short name	Operator	Prototype	Documentation	Supported types
andnb		Reg<T> andnb (const Reg<T> r1, const Reg<T> r1)	Computes the bitwise AND NOT: (~r1) & r2.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
andnb		Msk<N> andnb (const Msk<N> m1, const Msk<N> m2)	Computes the bitwise AND NOT: (~m1) & m2.	
orb	pipe and pipe=	Reg<T> orb (const Reg<T> r1, const Reg<T> r2)	Computes the bitwise OR: r1 pipe r2.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
orb	pipe and pipe=	Msk<N> orb (const Msk<N> m1, const Msk<N> m2)	Computes the bitwise OR: m1 pipe m2.	
xorb	^ and ^=	Reg<T> xorb (const Reg<T> r1, const Reg<T> r2)	Computes the bitwise XOR: r1 ^ r2.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
xorb	^ and ^=	Msk<N> xorb (const Msk<N> m1, const Msk<N> m2)	Computes the bitwise XOR: m1 ^ m2.	
lshift	<< and <<=	Reg<T> lshift (const Reg<T> r, const uint32_t n)	Computes the bitwise LEFT SHIFT: r << n.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
lshiftr	<< and <<=	Reg<T> lshiftr (const Reg<T> r1, const Reg<T> r2)	Computes the bitwise LEFT SHIFT: r1 << r2.	int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
lshift	<< and <<=	Msk<N> lshift (const Msk<N> m, const uint32_t n)	Computes the bitwise LEFT SHIFT: m << n.	
rshift	>> and >>=	Reg<T> rshift (const Reg<T> r, const uint32_t n)	Computes the bitwise RIGHT SHIFT: r >> n.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
rshiftr	>> and >>=	Reg<T> rshiftr (const Reg<T> r1, const Reg<T> r2)	Computes the bitwise RIGHT SHIFT: r1 >> r2.	int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t

Short name	Operator	Prototype	Documentation	Supported types
rshift	>> and >>=	Msk<N> rshift (const Msk<N> m, const uint32_t n)	Computes the bitwise RIGHT SHIFT: <code>m >> n</code> .	
notb	~	Reg<T> notb (const Reg<T> r)	Computes the bitwise NOT: <code>~r</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
notb	~	Msk<N> notb (const Msk<N> m)	Computes the bitwise NOT: <code>~m</code> .	

Logical Comparisons

Short name	Operator	Prototype	Documentation	Supported types
cmpeq	==	Msk<N> cmpeq (const Reg<T> r1, const Reg<T> r2)	Compares if equal to: <code>r1 == r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
cmpneq	!=	Msk<N> cmpneq (const Reg<T> r1, const Reg<T> r2)	Compares if not equal to: <code>r1 != r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
cmpge	>=	Msk<N> cmpge (const Reg<T> r1, const Reg<T> r2)	Compares if greater or equal to: <code>r1 >= r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
cmpgt	>	Msk<N> cmpgt (const Reg<T> r1, const Reg<T> r2)	Compares if strictly greater than: <code>r1 > r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
cmple	<=	Msk<N> cmple (const Reg<T> r1, const Reg<T> r2)	Compares if lower or equal to: <code>r1 <= r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
cmplt	<	Msk<N> cmplt (const Reg<T> r1, const Reg<T> r2)	Compares if strictly lower than: <code>r1 < r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t

Conversions and Packing

Short name	Prototype	Documentation	Supported types
------------	-----------	---------------	-----------------

Short name	Prototype	Documentation	Supported types
toReg	Reg<T> toReg (const Msk<N> m)	Converts the mask m into a register of type T , the number of elements N has to be the same for the mask and the register. If the mask is false then all the bits of the corresponding element are set to 0, otherwise if the mask is true then all the bits are set to 1 (be careful, for float datatypes true is interpreted as NaN!).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
cvt	Reg<T2> cvt (const Reg<T1> r)	Converts the elements of r into an other representation (the new representation and the original one have to have the same size).	float -> int32_t, float -> uint32_t, int32_t -> float, uint32_t -> float, double -> int64_t, double -> uint64_t, int64_t -> double, uint64_t -> double
cvt	Reg<T2> cvt (const Reg_2<T1> r)	Converts elements of r into bigger elements (in bits).	int8_t -> int16_t, uint8_t -> uint16_t, int16_t -> int32_t, uint16_t -> uint32_t, int32_t -> int64_t, uint32_t -> uint64_t
pack	Reg<T2> pack (const Reg<T1> r1, const Reg<T1> r2)	Packs elements of r1 and r2 into smaller elements (some information can be lost in the conversion).	int32_t -> int16_t, uint32_t -> uint16_t, int16_t -> int8_t, uint16_t -> uint8_t

Arithmetic Operations

Short name	Operator	Prototype	Documentation	Supported types
add	+ and +=	Reg<T> add (const Reg<T> r1, const Reg<T> r2)	Performs the arithmetic addition: r1 + r2 .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t

Short name	Operator	Prototype	Documentation	Supported types
sub	- and -=	Reg<T> sub (const Reg<T> r1, const Reg<T> r2)	Performs the arithmetic subtraction: $r1 - r2$.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
mul	* and *=	Reg<T> mul (const Reg<T> r1, const Reg<T> r2)	Performs the arithmetic multiplication: $r1 * r2$.	double, float, int32_t, int16_t, int8_t
div	/ and /=	Reg<T> div (const Reg<T> r1, const Reg<T> r2)	Performs the arithmetic division: $r1 / r2$.	double, float
fmadd		Reg<T> fmadd (const Reg<T> r1, const Reg<T> r2, const Reg<T> r3)	Performs the fused multiplication and addition: $r1 * r2 + r3$.	double, float
fnmadd		Reg<T> fnmadd (const Reg<T> r1, const Reg<T> r2, const Reg<T> r3)	Performs the negative fused multiplication and addition: $-(r1 * r2) + r3$.	double, float
fmsub		Reg<T> fmsub (const Reg<T> r1, const Reg<T> r2, const Reg<T> r3)	Performs the fused multiplication and subtraction: $r1 * r2 - r3$.	double, float
fnmsub		Reg<T> fnmsub (const Reg<T> r1, const Reg<T> r2, const Reg<T> r3)	Performs the negative fused multiplication and subtraction: $-(r1 * r2) - r3$.	double, float
min		Reg<T> min (const Reg<T> r1, const Reg<T> r2)	Selects the minimum: $r1_i < r2_i ? r1_i : r2_i$.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t

Short name	Operator	Prototype	Documentation	Supported types
max		Reg<T> max (const Reg<T> r1, const Reg<T> r2)	Selects the maximum: $r1_i > r2_i ? r1_i : r2_i$.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
div2		Reg<T> div2 (const Reg<T> r)	Performs the arithmetic division by two: $r / 2$.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
div4		Reg<T> div4 (const Reg<T> r)	Performs the arithmetic division by four: $r / 4$.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
abs		Reg<T> abs (const Reg<T> r)	Computes the absolute value of r .	double, float, int64_t, int32_t, int16_t, int8_t
sqrt		Reg<T> sqrt (const Reg<T> r)	Computes the square root of r .	double, float
rsqrt		Reg<T> rsqrt (const Reg<T> r)	Computes the reciprocal square root of r : $1 / \text{sqrt}(r)$.	double, float
sat		Reg<T> sat (const Reg<T> r, const T minv, const T maxv)	Saturates the register values: $\text{max}(\text{min}(r, \text{minv}), \text{maxv})$.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
neg		Reg<T> neg (const Reg<T> r, const Msk<N> m)	Negates the register elements following the mask values: $m_i ? -r_i : r_i$.	double, float, int64_t, int32_t, int16_t, int8_t
neg		Reg<T> neg (const Reg<T> r1, const Reg<T> r2)	Negates the register elements following the last register values: $r2_i < 0 ? -r1_i : r1_i$.	double, float, int64_t, int32_t, int16_t, int8_t
sign		Msk<N> sign (const Reg<T> r)	Returns the sign: $r < 0$.	double, float, int64_t, int32_t, int16_t, int8_t

Short name	Operator	Prototype	Documentation	Supported types
round		Reg<T> round (const Reg<T> r)	Rounds the register values: fractional_part(r) >= 0.5 ? integral_part(r) + 1 : integral_part(r).	double, float
trunc		Reg<T> trunc (const Reg<T> r)	Truncates the register values: integral_part(r) .	double, float

Arithmetic Operations on Complex Numbers

The complex operations are exclusively performed on `Regx2<T>` objects (one `Regx2<T>` object contains two `Reg<T>` hardware registers). Each `Regx2<T>` object contains `mipp::N<T>()` complex number. If we declare a `Regx2<T> cmplx` object, the `cmplx[0]` register will contain the real part of the complex numbers and `cmplx[1]` will contain the imaginary part. Depending on how you stored your complex numbers in memory you can need to use reordering before calling a complex operation. For instance, if you choose to store the complex numbers in a mixed format like this: `r0, i0, r1, i1, r2, i2, ..., rn, in` you will need to call the `mipp::deinterleave` operation before and the `mipp::interleave` operation after the complex operation.

Short name	Operator	Prototype	Documentation	Supported types
cadd	+ and +=	Regx2<T> cadd (const Regx2<T> r1, const Regx2<T> r2)	Performs the complex addition: <code>r1</code> + <code>r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
csub	- and -=	Regx2<T> csub (const Regx2<T> r1, const Regx2<T> r2)	Performs the complex subtraction: <code>r1</code> - <code>r2</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
cmul	* and *=	Regx2<T> cmul (const Regx2<T> r1, const Regx2<T> r2)	Performs the complex multiplication: <code>r1</code> * <code>r2</code> .	double, float, int32_t, int16_t, int8_t
cdiv	/ and /=	Regx2<T> cdiv (const Regx2<T> r1, const Regx2<T> r2)	Performs the complex division: <code>r1</code> / <code>r2</code> .	double, float
cmulconj		Regx2<T> cmulconj (const Regx2<T> r1, const Regx2<T> r2)	Performs the complex multiplication with conjugate: <code>r1</code> * <code>conj(r2)</code> .	double, float, int32_t, int16_t, int8_t

Short name	Operator	Prototype	Documentation	Supported types
conj		Regx2<T> cmulconj (const Regx2<T> r)	Computes the conjugate: conj(r).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
norm		Reg <T> norm (const Regx2<T> r)	Computes the squared magnitude: norm(r).	double, float, int32_t, int16_t, int8_t

Reductions (Horizontal Functions)

Short name	Operator	Prototype	Documentation	Supported types
hadd or sum		T hadd (const Reg<T> r)	Sums all the elements in the register r: r_1 + r_2 + ... + r_n.	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
hmul		T hmul (const Reg<T> r)	Multiplies all the elements in the register r: r_1 * r_2 * ... * r_n.	double, float, int64_t, int32_t, int16_t, int8_t
hmin		T hmin (const Reg<T> r)	Selects the minimum element in the register r: min(min(min(..., r_1), r_2), r_n).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
hmax		T hmax (const Reg<T> r)	Selects the maximum element in the register r: max(max(max(..., r_1), r_2), r_n).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
testz		bool testz (const Reg<T> r1, const Reg<T> r2)	Mainly tests if all the elements of the registers are zeros: r = (r1 & r2); !(r_1 OR r_2 OR ... OR r_n).	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
testz		bool testz (const Msk<N> m1, const Msk<N> m2)	Mainly tests if all the elements of the masks are zeros: m = (m1 & m2); !(m_1 OR m_2 OR ... OR m_n).	

Short name	Prototype	Documentation	Supported types
testz	bool testz (const Reg<T> r)	Tests if all the elements of the register are zeros: <code>!(r_1 OR r_2 OR ... OR r_n)</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t
testz	bool testz (const Msk<N> m)	Tests if all the elements of the mask are zeros: <code>!(m_1 OR m_2 OR ... OR m_n)</code> .	
Reduction<T,OP>	T Reduction<T,OP>::apply (const Reg<T> r)	Generic reduction operation, can take a user defined operator <code>OP</code> and will performs the reduction with it on <code>r</code> .	double, float, int64_t, uint64_t, int32_t, uint32_t, int16_t, uint16_t, int8_t, uint8_t

Math Functions

Short name	Prototype	Documentation	Supported types
exp	Reg<T> exp (const Reg<T> r)	Computes the exponential of <code>r</code> .	double (only on icpc), float
log	Reg<T> log (const Reg<T> r)	Computes the logarithm of <code>r</code> .	double (only on icpc), float
sin	Reg<T> sin (const Reg<T> r)	Computes the sines of <code>r</code> .	double (only on icpc), float
cos	Reg<T> cos (const Reg<T> r)	Computes the cosines of <code>r</code> .	double (only on icpc), float
tan	Reg<T> tan (const Reg<T> r)	Computes the tangent of <code>r</code> .	double (only on icpc), float
sincos	void sincos (const Reg<T> r, Reg<T>& s, Reg<T>& c)	Computes at once the sines (in <code>s</code>) and the cosines (in <code>c</code>) of <code>r</code> .	double (only on icpc), float
sincos	Regx2<T> sincos (const Reg<T> r)	Computes and returns at once the sines and the cosines of <code>r</code> .	double (only on icpc), float
cossin	Regx2<T> cossin (const Reg<T> r)	Computes and returns at once the cosines and the sines of <code>r</code> .	double (only on icpc), float
sinh	Reg<T> sinh (const Reg<T> r)	Computes the hyperbolic sines of <code>r</code> .	double (only on icpc), float

Short name	Prototype	Documentation	Supported types
cosh	Reg<T> cosh (const Reg<T> r)	Computes the hyperbolic cosines of <i>r</i> .	double (only on icpc), float
tanh	Reg<T> tanh (const Reg<T> r)	Computes the hyperbolic tangent of <i>r</i> .	double (only on icpc), float
asinh	Reg<T> asinh (const Reg<T> r)	Computes the inverse hyperbolic sines of <i>r</i> .	double (only on icpc), float
acosh	Reg<T> acosh (const Reg<T> r)	Computes the inverse hyperbolic cosines of <i>r</i> .	double (only on icpc), float
atanh	Reg<T> atanh (const Reg<T> r)	Computes the inverse hyperbolic tangent of <i>r</i> .	double (only on icpc), float

ARM SVE

SVE Length Specific

An ARM SVE version is under construction. This version uses *SVE length specific* which is more appropriated to the MIPP architecture. This way, the size of the *MIPP registers* is defined at the compilation time. As a reminder, the vector length can vary from a minimum of 128 bits up to a maximum of 2048 bits, at 128-bit increments. On GNU and Clang compilers, it is specified at the compilation time with the `-msve-vector-bits=<size>` flag.

Supported MIPP Operations

- **Memory operations:** `load`, `store`, `blend`, `set`, `set1`, `gather`, `scatter`, `maskzld`, `maskst`, `maskzgat`, `maskzca`
- **Logical comparisons:** `cmpeq`, `cmneq`
- **Bitwise operations:** `andb`, `notb` (*msk*)
- **Arithmetic operations:** `fmadd`, `add`, `sub`, `mul`, `div`
- **Reductions:** `testz` (*msk*), `Reduce<T>`, `add`

Byte and *word* operations are not yet implemented.

How to cite MIPP

We recommend you to cite the following article:

- Adrien Cassagne, Olivier Aumage, Denis Barthou, Camille Leroux and Christophe Jégo, **MIPP: a Portable C++ SIMD Wrapper and its use for Error Correction Coding in 5G Standard**, *The 5th International Workshop on Programming Models for SIMD/Vector Processing (WPMVP 2018), February 2018.*